

Actualizing Ethical Principles for Curating Large-Scale Training Datasets in the Era of Massive AI Models

Silvia Cazacu¹, Alice Qian², Dora Zhao³, Kathleen Pine⁴, Shawn Walker⁴, Hong Shen², Laura Dabbish², Georgia Panagiotidou⁵, Morgan Klaus Scheuerman⁶

¹Télécom Paris, Institut Polytechnique de Paris, France ²Carnegie Mellon University, Pittsburgh, USA

³Stanford University, Stanford, USA ⁴Arizona State University, Tempe, USA

⁵King's College London, London, UK ⁶Sony AI, Barcelona, Spain

Abstract

While AI technologies are often framed as ubiquitous and inevitable, their expansion relies on the large-scale extraction of data from diverse global communities. However, the datasets powering foundation models are often treated as found artifacts rather than products of specific power dynamics and human labor. Current practices frequently disregard the structural inequities embedded in data, even as these systems profoundly impact systemically marginalized communities. While frameworks for ethical curation exist for smaller datasets, the massive scale of foundation models has introduced a logic of extraction that prioritizes volume over accountability. This workshop invites researchers, practitioners, and activists to move beyond standard technical hurdles and instead problematize the foundational assumptions of large-scale data work. This workshop builds on a series of ongoing conversations across scholarly communities and continues work initiated at CSCW 2025. Drawing from the CRAFT tradition of transdisciplinary exchange and community action, we will facilitate a collective refactoring of the three core pillars of data curation: composition, focused on interrogating whose lives are extracted and how representation is shaped by hegemonic interests; process, which centers the invisible labor and situated contexts involved in curating and cleaning massive data stores, and release, focused on rethinking the governance, accountability, and potential for refusal in how these models are shared with the world. Our goal is to cultivate a community-led conceptual framework that reimagines more just sociotechnical futures—transforming data curation from a top-down technical requirement into an act of collective responsibility and repair.

Keywords

Fairness, ethics, responsible AI, foundation models, generative AI, datasets, machine learning, responsible artificial intelligence, human-centric artificial intelligence, algorithmic bias, values in design, work practice

This is a workshop paper submitted and accepted to the CRAFT track at the ACM Conference on Fairness, Accountability, and Transparency (FAccT 2026). This is a non-archival manuscript and is not part of the official proceedings.

1 Description

The unfettered expansion of AI infrastructure relies on the systematic erasure of the power dynamics and labor exploitation inherent in dataset production. While small-scale ethical curation frameworks exist, the move toward massive foundation models has institutionalized an extractive logic that treats data as a boundless

resource and prioritizes computational volume over community accountability. This CRAFT aims to rethink the premises of boundless data curation while opening new paths for sociotechnical justice and repair.

The CRAFT will bring together community activists, researchers, and practitioners to discuss scenarios of data curation regarding the composition of the data itself, the process of curation, and the eventual release of datasets for AI training. More specifically, based on real-world scenarios presented by community activists and practitioners, participants will: 1) engage in a scenario-based role-play, adopting the viewpoints of developers, policy advocates, or impacted community members to problematize the systemic limitations inherent in each position; and 2) use these insights to collectively reimagine more ethical and just possibilities through speculative futures that respond directly to the real-world harms identified.

Responsible dataset curation is not simple or straightforward across all different stages of the dataset lifecycle, from composition to curation to release. Even in instances where researchers have attempted to collect fair or otherwise ethical datasets, they have been found to incidentally violate the expectations of other researchers or the public [14, 45, 49]. Existing datasets have violated data subject consent [3, 6, 29, 33], infringed on copyright [12, 20, 32], exploited data workers [4, 16, 44], contained poor demographic distribution [8, 18, 47] and offensive image labels [5, 7, 40], and accidentally included illegal content [39]. It seems that collecting datasets responsibly is extremely challenging, even for relatively small datasets. Recent work by Zhao et al. uncovered extensive challenges that dataset curators face in enacting ethics throughout the dataset curation lifecycle [48]. The challenges that arise when creating fair evaluation datasets may be much more complex for large-scale training datasets — or even entirely different.

While there will undoubtedly be outstanding challenges in responsibly curating large-scale datasets that adhere to ethical principles, it remains necessary to address them in our shared goal of creating ethical foundation models downstream. Scholars have already made numerous suggestions for positively improving dataset qualities like increasing data diversity [24, 26], obtaining data subject consent [22, 23, 43], providing fair wages to data workers [1, 13, 36], limiting dataset use [28], engaging stakeholders in data taxonomy design [25, 27, 35], and creating transparent documentation [11, 19, 46]. Andrews et al. provided a comprehensive framework of ethical considerations for responsible dataset curation throughout the development lifecycle, illuminating idealistic data curation approaches for smaller scale evaluation datasets [2].

Yet, many of these ethical approaches may need to be reconsidered and redefined for the scale of foundation model data. This creates the urgency to reimagine entirely new approaches which have yet to be considered, especially around issues like data instability and recency, transparency tools for parsing massive unstructured datasets, and environmental stability for both collection and use [17, 37, 42].

Local activists have already identified opportunities and developed resistance tactics. The Kaitiākitanga License, originating with the Māori people of New Zealand, grants use of indigenous language data for training foundation models only if it benefits the community while offering community ‘guardianship’ (Kaitiākitanga) over the data in perpetuity [38]. Mukurtu is a community-driven platform that empowers indigenous groups to manage their digital heritage through granular access protocols, where some information is visible only to certain families, elders, or genders [10]. In Canada, the OCAP (Ownership, Control, Access, and Possession) principles for data governance ensure that First Nations communities have the right to own and control data curation processes that affect them [15].

However, all these perspectives and tactics cannot be mobilized efficiently and meaningfully without deeper engagement across communities.

2 Goal and Collective Outcomes

With growing interest in foundation model implications across HCI and responsible AI communities(e.g., [9, 21, 41]), this CRAFT brings together community activists, researchers, policymakers, and practitioners to move from problem identification to framework development. Building on a successful workshop at CSCW 2025 [34]¹, where participants identified key tensions and trade-offs across these areas. Now at FAccT, *we move from exploration to framework development*. Through scenario analysis and speculative futures exercises, participants will synthesize these insights into actionable principles that can guide responsible dataset curation at scale. This CRAFT aims to bring together community activists, researchers, policymakers, and practitioners to reflect on and reimagine how the process of dataset curation for foundation models might be supported ethically.

3 Activities

Our CRAFT will be organized as a 3-hour workshop in hybrid format or in-person format. We have planned the activities to be flexible to a hybrid format to be inclusive to potentially interested practitioners who may not be able to travel. The CRAFT has two main activities: (1) a round-table discussion and (2) a speculative futures exercise in groups, prefaced by a plenary introduction and followed by plenary closing remarks. The purpose of the discussion is to foster collaboration among participants, starting from a set of real-world problems, and to encourage them to discuss their own experiences and needs. The speculative exercise invites participants to envision an ideal form of ethical data curation and collaboratively

develop an imaginary AI workflow. We describe session activities as follows:

Introduction (30 min): The session begins with opening plenary remarks from the organizers, introducing participants to the agenda of the day.

Session 1, Round-Table Discussion (45 min): The round table discussion will be supported by scenarios for dataset curation. A wide set of scenarios will be developed before the CRAFT by the organization team based on their (local and personal) knowledge and experience with dataset curation, and they will focus on three aspects of the data curation process: the composition of the data itself, the process of curation, and the eventual release of datasets for AI training. Each scenario will contain a description of a dataset curation context; a conflict, critique or best practice example related to composition, curation or process; and a set of provocative questions to start a group reflection. Participants will have time to familiarize themselves with the scenarios, after which the CRAFT facilitators, from the organization team, will lead the discussion, focusing on one scenario at a time. Three scenarios will be discussed in total, each scenario corresponding to either data curation composition, process, release aspects. No prior knowledge of the scenarios will be assumed/expected, and the discussion, depending on the attendance size, will happen as a plenary session. If the group attending is too large for one plenary discussion, we will set up each scenario at a different station and invite participants to rotate through the stations. At each station, they find a scenario and a provocation board.

[Coffee break 1 (15 min)]

Session 2, Speculative Futures Exercise (60 min): The second part separates participants into groups where they are guided to reimagine these scenarios by role-playing different perspectives, adopting the viewpoints of developers, policy advocates, or community members, guided by instructions. Here, participants will transition from existing limitations to possible futures. To help their imagination, each group will receive a constraint card along with role-play instructions. Based on these instructions, participants will create group posters, which they will present to the plenary, prompting a speculative reflection between all groups, facilitated by the organization team. This reflection will help participants to collectively reimagine more ethical and just possibilities through speculative futures that respond directly to the real-world harms identified.

[Coffee break 2 (10 min)]

Closing remarks (20 min): organizers will give closing remarks and provide opportunities for participants to opt into follow-up efforts from the session and stay in contact with each other.

4 Artifacts

After the conclusion of the CRAFT, the organizers will provide a summary of the CRAFT on the website² and Discord channel or Slack space. As this is part of a series of workshops that was started at CSCW in 2025, we aim to continue organizing events focused on ideating and refining best practices for ethically curating large-scale training datasets with different scholarly communities, including: at machine learning conferences (NeurIPS), design conferences (DIS),

¹The website of the previous workshop at CSCW <https://responsible-data-workshop.github.io/cscw2025/> and the workshop tools repository <https://responsiblefndata.github.io/>

²<https://responsiblefndata.github.io/>

HCI conferences (CHI, CSCW) and responsible AI conferences (AIES).

We plan to use Discord or Slack for future workshops in the workshop series. This will facilitate continued participation and community building as the workshop series develops across conference communities. We will then collate and analyze our longitudinal takeaways in the form of an article or white paper so that the broader community can learn from the shared knowledge of the workshop community. Interested CRAFT participants will be invited to contribute to this article or white paper. We also hope to publish a special issue of a journal focused on ethical principles, challenges, and opportunities for dataset curation. This special issue will serve as a platform for workshop participants to either expand upon their position papers, refining them for potential publication after the workshop, or to submit any other pertinent work they may have developed.

5 Audience

We aim to bring together researchers studying dataset ethics and foundation models, practitioners building AI systems at scale, community activists working on data sovereignty and digital rights, and policymakers developing AI governance frameworks. We expect around 15-35 total participants. We plan to cap attendance to 35 to facilitate more engaged conversations in the limited time of the CRAFT.

6 Organizers and Facilitators

Silvia Cazacu is a post-doctoral researcher at Télécom Paris, Institut Polytechnique de Paris, where she works on the intersection between democratic deliberation and data physicalization. Her research applies a feminist lens to critically analyze the power dynamics that emerge between diverse stakeholders who collaborate on infrastructural decisions about data. She combines participatory methods with tangible thinking tools to make complex issues visible and negotiable.

Alice Qian is a PhD student at Carnegie Mellon University's Human-Computer Interaction Institute, where her research focuses on examining methods of leveraging and amplifying the human infrastructure of responsible AI development. She has specific experience studying AI safety practices within organizational settings [30] as well as through crowdsourcing efforts [31]. She has co-organized two previous CSCW workshops focusing on AI red teaming. She also has experience red teaming in practice as well as in the analysis of results.

Dora Zhao is a PhD student at Stanford University. Her research focuses on understanding how we can design AI systems that consider long-range benefits for users, rather than optimizing only for short-term gains. Her work includes measuring emerging harms of AI systems and building new technical methods towards mitigating these harms.

Kathleen Pine is an Associate Professor in the College of Health Solutions at Arizona State University. Her research examines design, implementation, and use of health information technologies, and has researched how increased demands for data work are re-configuring healthcare practice.

Shawn Walker is an Assistant Professor in the School of Social and Behavioral Sciences at Arizona State University. His research focuses on mis and disinformation, and the challenges of collecting, analyzing, and preserving data from social media platforms. He also examines the use and implications of social media and web archives to train ML models.

Hong Shen is an Assistant Professor in the Human-Computer Interaction Institute at Carnegie Mellon, where she directs the CARE (Collective AI Research and Evaluation) lab. She has received a Google Research Scholar Award, a Microsoft AI and Society Fellowship, and Best Paper Awards at top HCI venues.

Laura Dabbish is a Professor at Carnegie Mellon University in the Human-Computer Interaction Institute within the School of Computer Science. She directs the Connected Experience Lab, which addresses social, technical, and policy issues in emerging technology.

Georgia Panagiotidou is an Assistant Professor at King's College London at the department of Informatics. Her work broadly centers around how people engage with environmental data exploring key challenges such as biases, uncertainties, and friction in collective data use. Georgia investigates the environmental impact of AI commencing from the tools and interfaces that developers, end-users and communities experience as probes with the goal to build bridges between the domains of HCI, STS and environmentally responsible AI.

Morgan Klaus Scheuerman is a research scientist on the AI Ethics team at Sony AI, where he currently works on the ethical implications of generative AI for creative industries.

Acknowledgments

The work of Silvia Cazacu was supported by a French government grant managed by the Agence Nationale de la Recherche as part of the France 2030 program, reference ANR-22-EXEN-0005 (PEPR eNSEMBLE / PC4 CONGRATS). Dora Zhao is partially supported by the P.D. Soros Fellowship for New Americans. Alice Qian was supported by the National Institute of Standards and Technology (NIST) under Federal Award ID Number 60NANB24D231 and the Carnegie Mellon University³ AI Measurement Science and Engineering Center (AIMSEC).

References

- [1] Antonio Aloisi. 2015. Commoditized Workers: Case Study Research on Labor Law Issues Arising from a Set of on-Demand/Gig Economy Platforms. *Comparative Labor Law & Policy Journal* 37 (2015).
- [2] Jerone Andrews, Dora Zhao, William Thong, Apostolos Modas, Orestis Papakyriakopoulos, and Alice Xiang. 2023. Ethical Considerations for Responsible Data Curation. In *Thirty-Seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*.
- [3] Jerone TA Andrews, Dora Zhao, William Thong, Apostolos Modas, Orestis Papakyriakopoulos, Shruti Nagpal, and Alice Xiang. 2023. Ethical considerations for collecting human-centric image datasets. *arXiv preprint arXiv:2302.03629* (2023).
- [4] Jo Bates, Elli Gerakopoulou, and Alessandro Checco. 2023. Addressing Labour Exploitation in the Data Science Pipeline: Views of Precarious US-based Crowdworkers on Adversarial and Co-Operative Interventions. *Journal of Information, Communication and Ethics in Society* 21, 3 (Jan. 2023), 342–357. <https://doi.org/10.1108/JICES-08-2022-0069>
- [5] Abeba Birhane, Sepehr Dehdashtian, Vinay Prabhu, and Vishnu Boddeti. 2024. The Dark Side of Dataset Scaling: Evaluating Racial Classification in Multimodal Models. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability,*

³<https://ror.org/05x2bcf33>

- and Transparency (FAccT '24). Association for Computing Machinery, New York, NY, USA, 1229–1244. <https://doi.org/10.1145/3630106.3658968>
- [6] Abeba Birhane and Vinay Uday Prabhu. 2021. Large image datasets: A pyrrhic win for computer vision?. In *IEEE Winter Conference on Applications of Computer Vision (WACV)*. 1536–1546.
 - [7] Abeba Birhane and Vinay Uday Prabhu. 2021. Large Image Datasets: A Pyrrhic Win for Computer Vision?. In *2021 IEEE Winter Conference on Applications of Computer Vision (WACV)*. 1536–1546. <https://doi.org/10.1109/WACV48630.2021.00158>
 - [8] Alceu Bissoto, Catarina Barata, Eduardo Valle, and Sandra Avila. 2024. Even Small Correlation and Diversity Shifts Pose Dataset-Bias Issues. *Pattern Recognition Letters* 179 (March 2024), 87–93. <https://doi.org/10.1016/j.patrec.2024.01.026>
 - [9] Youjin Choi, JaeYoung Moon, Kyung-Joong Kim, and Jin-Hyuk Hong. 2024. Exploring the Potential of Generative AI in Song-Signing. In *Companion of the 2024 on ACM International Joint Conference on Pervasive and Ubiquitous Computing (UbiComp '24)*. Association for Computing Machinery, New York, NY, USA, 816–820. <https://doi.org/10.1145/3675094.3678378>
 - [10] Kimberly Christen, Michael Wynne, and Lotus Norton-Wisla. 2024. Mukurtu CMS: An Indigenous Archive and Publishing Tool. <https://mukurtu.org/Version4.0>
 - [11] Roxana Daneshjou, Mary P. Smith, Mary D. Sun, Veronica Rotemberg, and James Zou. 2021. Lack of Transparency and Potential Bias in Artificial Intelligence Data Sets and Algorithms: A Scoping Review. *JAMA Dermatology* 157, 11 (Nov. 2021), 1362–1369. <https://doi.org/10.1001/jamadermatol.2021.3129>
 - [12] Devansh. 2023. Data Laundering: How Stability AI Managed to Get Millions of Copyrighted Artworks without Paying....
 - [13] Mark Díaz, Ian Kivlichan, Rachel Rosen, Dylan Baker, Razvan Amironesei, Vinodkumar Prabhakaran, and Emily Denton. 2022. CrowdWorkSheets: Accounting for Individual and Collective Identities Underlying Crowdsourced Dataset Annotation. In *2022 ACM Conference on Fairness, Accountability, and Transparency*. Association for Computing Machinery, New York, NY, USA, 2342–2351. <https://doi.org/10.1145/3531146.3534647>
 - [14] Vittoria Elliott. 2024. AI Tools Are Secretly Training on Real Images of Children. *Wired* (June 2024).
 - [15] First Nations Information Governance Centre. 2020. The First Nations Principles of OCAP®. <https://fnigc.ca/ocap-training/>. Accessed: 2026-03-23.
 - [16] Mary L. Gray and Suri Siddharth. 2019. *Ghost Work: How to Stop Silicon Valley from Building a New Global Underclass*.
 - [17] Philipp Hacker. 2024. Sustainable AI Regulation. *Common Market Law Review* 61, 2 (April 2024).
 - [18] Yusuke Hirota, Yuta Nakashima, and Noa Garcia. 2022. Gender and Racial Bias in Visual Question Answering Datasets. In *2022 ACM Conference on Fairness, Accountability, and Transparency*. 1280–1292. <https://doi.org/10.1145/3531146.3533184> arXiv:2205.08148 [cs]
 - [19] Oana Inel, Tim Draws, and Lora Aroyo. 2023. Collect, Measure, Repeat: Reliability Factors for Responsible AI Data Collection. *Proceedings of the AAAI Conference on Human Computation and Crowdsourcing* 11, 1 (Nov. 2023), 51–64. <https://doi.org/10.1609/hcomp.v11i1.27547>
 - [20] Yotam Kaplan and Ayelet Gordon-Tapiro. 2024. Generative AI Training as Unjust Enrichment.
 - [21] Chinmay Kulkarni, Tongshuang Wu, Kenneth Holstein, Q. Vera Liao, Min Kyung Lee, Mina Lee, and Hariharan Subramonyam. 2023. LLMs and the Infrastructure of CSCW. In *Companion Publication of the 2023 Conference on Computer Supported Cooperative Work and Social Computing (CSCW '23 Companion)*. Association for Computing Machinery, New York, NY, USA, 408–410. <https://doi.org/10.1145/3584931.3608438>
 - [22] Shayne Longpre, Robert Mahari, Ariel Lee, Campbell Lund, Hamidah Oderinwale, William Brannon, Nayan Saxena, Naana Obeng-Marnu, Tobin South, Cole Hunter, Christopher Klamm, Hailey Schoelkopf, Nikhil Singh, Manuel Cherep, Mustafa Anis, An Dinh, Caroline Chitongo, Da Yin, Damien Sileo, Deividas Mataciunas, Diganta Misra, Emad Alghamdi, Enrico Shippole, Jianguo Zhang, Joanna Materzynska, Kun Qian, Kush Tiwary, Lester Miranda, Manan Dey, Minnie Liang, Niklas Muennighoff, Seonghyeon Ye, Seungone Kim, Shrestha Mohanty, Vivek Sharma, Vu Minh Chien, Xuhui Zhou, Yizhi Li, Caiming Xiong, Luis Villa, Stella Biderman, Hanlin Li, Daphne Ippolito, Sara Hooker, and Jad Kabbara. 2024. Consent in Crisis: The Rapid Decline of the AI Data Commons. (2024).
 - [23] Kate McCandless. 2021. Just Because the Data Is There, It Doesn't Mean It's Yours to Take : Exploring User, Researcher and Review Board Perceptions in Twitter Data Research. *Emerging Library & Information Perspectives* 4, 1 (July 2021), 34–61. <https://doi.org/10.5206/elp.v4i1.13554>
 - [24] Keziah Naggita, Julianne LaChance, and Alice Xiang. 2023. Flickr Africa: Examining Geo-Diversity in Large-Scale, Human-Centric Visual Data. In *Proceedings of the 2023 AAAI/ACM Conference on AI, Ethics, and Society (AI/ES '23)*. Association for Computing Machinery, New York, NY, USA, 520–530. <https://doi.org/10.1145/3600211.3604659>
 - [25] Cecilia Panigutti, Andrea Beretta, Daniele Fadda, Fosca Giannotti, Dino Pedreschi, Alan Perotti, and Salvatore Rinzivillo. 2023. Co-Design of Human-centered, Explainable AI for Clinical Decision Support. *ACM Trans. Interact. Intell. Syst.* 13, 4 (Dec. 2023), 21:1–21:35. <https://doi.org/10.1145/3587271>
 - [26] Orestis Papakyriakopoulos, Anna Seo Gyeong Choi, William Thong, Dora Zhao, Jerone Andrews, Rebecca Bourke, Alice Xiang, and Allison Konecke. 2023. Augmented Datasheets for Speech Datasets and Ethical Decision-Making. In *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency (FAccT '23)*. Association for Computing Machinery, New York, NY, USA, 881–904. <https://doi.org/10.1145/3593013.3594049>
 - [27] Addie Payne Morgan, Bryan F. Howell, and Grace Kilbourn-Barber. 2024. CO-DESIGN AND ARTIFICIAL INTELLIGENCE: A METHOD TO EMPOWER END-USERS IN VISUAL COMMUNICATION. In *DS 131: Proceedings of the International Conference on Engineering and Product Design Education (E&PDE 2024)*. 557–562. <https://doi.org/10.35199/EPDE.2024.94>
 - [28] Kenny Peng, Arunesh Mathur, and Arvind Narayanan. 2021. Mitigating Dataset Harms Requires Stewardship: Lessons from 1000 Papers. (Aug. 2021). arXiv:2108.02922
 - [29] Kenneth L Peng, Arunesh Mathur, and Arvind Narayanan. 2021. Mitigating dataset harms requires stewardship: Lessons from 1000 papers. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 2)*.
 - [30] Alice Qian, Judith Amores, Hong Shen, Mary Czerwinski, Mary L Gray, and Jina Suh. 2025. AURA: Amplifying Understanding, Resilience, and Awareness for Responsible AI Content Work. *Proceedings of the ACM on Human-Computer Interaction* 9, CSCW2 (2025), 1–45.
 - [31] Alice Qian, Ryland Shaw, Laura Dabbish, Jina Suh, and Hong Shen. 2025. Locating Risk: Task Designers and the Challenge of Risk Disclosure in RAI Content Work. arXiv preprint arXiv:2505.24246 (2025).
 - [32] Pamela Samuelson. 2023. Generative AI Meets Copyright. *Science* 381, 6654 (July 2023), 158–161. <https://doi.org/10.1126/science.ad0656>
 - [33] Morgan Klaus Scheuerman, Alex Hanna, and Emily Denton. 2021. Do datasets have politics? Disciplinary values in computer vision dataset development. *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW2 (2021), 1–37.
 - [34] Morgan Klaus Scheuerman, Dora Zhao, Jerone TA Andrews, Abeba Birhane, Q Vera Liao, Georgia Panagiotidou, Pooja Chitre, Kathleen Pine, Shawn Walker, Jieyu Zhao, et al. 2025. Responsibly Training Foundation Models: Actualizing Ethical Principles for Curating Large-Scale Training Datasets in the Era of Massive AI Models. In *Companion Publication of the 2025 Conference on Computer-Supported Cooperative Work and Social Computing*. 99–105.
 - [35] Cathrine Seidelin, Yvonne Dittrich, and Erik Grönvall. 2020. Foregrounding Data in Co-Design – An Exploration of How Data May Become an Object of Design. *International Journal of Human-Computer Studies* 143 (Nov. 2020), 102505. <https://doi.org/10.1016/j.ijhcs.2020.102505>
 - [36] Boaz Shmueli, Jan Fell, Soumya Ray, and Lun-Wei Ku. 2021. Beyond Fair Pay: Ethical Implications of NLP Crowdsourcing. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. 3758–3769. <https://doi.org/10.18653/v1/2021.naacl-main.295> arXiv:2104.10097 [cs]
 - [37] Emma Strubell, Ananya Ganesh, and Andrew McCallum. 2019. Energy and Policy Considerations for Deep Learning in NLP. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, Anna Korhonen, David Traum, and Lluís Màrquez (Eds.). Association for Computational Linguistics, Florence, Italy, 3645–3650. <https://doi.org/10.18653/v1/P19-1355>
 - [38] Te Reo Irirangi o Te Hiku o Te Ika. 2020. The Kaitiakitanga License. <https://kaitiaki.nz/kaitiakitanga> Indigenous data sovereignty and guardianship license.
 - [39] David Thiel. 2023. Investigation Finds AI Image Generation Models Trained on Child Abuse. <https://fsi.stanford.edu/news/investigation-finds-ai-image-generation-models-trained-child-abuse>.
 - [40] Eddie L. Ungless, Björn Ross, and Anne Lauscher. 2023. Stereotypes and Smut: The (Mis)Representation of Non-cisgender Identities by Text-to-Image Models. (May 2023). arXiv:2305.17072
 - [41] Priyan Vaithilingam, Ian Arawjo, and Elena L. Glassman. 2024. Imagining a Future of Designing with AI: Dynamic Grounding, Constructive Negotiation, and Sustainable Motivation. In *Proceedings of the 2024 ACM Designing Interactive Systems Conference (DIS '24)*. Association for Computing Machinery, New York, NY, USA, 289–300. <https://doi.org/10.1145/3643834.3661525>
 - [42] Hamish van der Ven, Diego Corry, Rawie Elnur, Viola Jasmine Provost, and Muh Syukron. 2024. Generative AI and Social Media May Exacerbate the Climate Crisis. *Global Environmental Politics* 24, 2 (May 2024), 9–18. https://doi.org/10.1162/glep_a_00747
 - [43] Lauren Wilcox, Robin Brewer, and Fernando Diaz. 2023. AI Consent Futures: A Case Study on Voice Data Collection with Clinicians. *Proc. ACM Hum.-Comput. Interact.* 7, CSCW2 (Oct. 2023), 316:1–316:30. <https://doi.org/10.1145/3610107>
 - [44] Adrienne Williams, Milagros Miceli, and Timnit Gebre. 2022. The Exploited Labor behind Artificial Intelligence. *NOEMA* (2022), 1–11.
 - [45] Chloe Xiang. 2023. OpenAI and Microsoft Sued for \$3 Billion Over Alleged ChatGPT 'Privacy Violations'.
 - [46] Meg Young, Luke Rodriguez, Emily Keller, Feiyang Sun, Boyang Sa, Jan Whittington, and Bill Howe. 2019. Beyond Open vs. Closed: Balancing Individual Privacy

- and Public Accountability in Data Sharing. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '19)*. Association for Computing Machinery, New York, NY, USA, 191–200. <https://doi.org/10.1145/3287560.3287577>
- [47] Dora Zhao, Jerone Andrews, Orestis Papakyriakopoulos, and Alice Xiang. 2024. Position: Measure Dataset Diversity, Don't Just Claim It. In *Proceedings of the 41st International Conference on Machine Learning*. PMLR, 60644–60673.
- [48] Dora Zhao, Morgan Klaus Scheuerman, Pooja Chitre, Jerone T. A. Andrews, Georgia Panagiotidou, Shawn Walker, Kathleen H. Pine, and Alice Xiang. 2024. A Taxonomy of Challenges to Curating Fair Datasets. In *NeurIPS 2024*. <https://doi.org/10.48550/arXiv.2406.06407> arXiv:2406.06407 [cs]
- [49] Michael Zimmer. 2010. "But the Data Is Already Public": On the Ethics of Research in Facebook. *Ethics and Information Technology* 12, 4 (Dec. 2010), 313–325. <https://doi.org/10.1007/s10676-010-9227-5>